

CREDIT CARD FRAUD DETECTION

DATA COLLECTION AND DATA PREPARATION

Import Libraries

```
In [1]: #import library
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sqlalchemy import create_engine
from sklearn.preprocessing import StandardScaler
```

Database connection and load data

```
In [2]: #create SQLAlchemy engine
engine = create_engine('postgresql+psycopg2://postgres:admin@localhost:5432/fraud_detectionDB')
```

✓ Deliverable: Data successfully stored in PostgreSQL, schema applied

```
In [3]: # read data into pandas
df = pd.read_sql_query("SELECT * FROM cc_data", engine)
```

```
In [4]: # View top rows
print(df.head())
```

	time	v1	v2	v3	v4	v5	v6	v7	\
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	

	v8	v9	...	v21	v22	v23	v24	v25	\
0	0.098698	0.363787	...	-0.018307	0.277838	-0.110474	0.066928	0.128539	
1	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288	-0.339846	0.167170	
2	0.247676	-1.514654	...	0.247998	0.771679	0.909412	-0.689281	-0.327642	
3	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321	-1.175575	0.647376	
4	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458	0.141267	-0.206010	

	v26	v27	v28	amount	class
0	-0.189115	0.133558	-0.021053	149.62	0
1	0.125895	-0.008983	0.014724	2.69	0
2	-0.139097	-0.055353	-0.059752	378.66	0
3	-0.221929	0.062723	0.061458	123.50	0
4	0.502292	0.219422	0.215153	69.99	0

[5 rows x 31 columns]

Data Profiling

```
In [5]: print("Dataset Information")
print(df.info())
```

Dataset Information

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 284807 entries, 0 to 284806

Data columns (total 31 columns):

#	Column	Non-Null	Count	Dtype
0	time	284807	non-null	float64
1	v1	284807	non-null	float64
2	v2	284807	non-null	float64
3	v3	284807	non-null	float64
4	v4	284807	non-null	float64
5	v5	284807	non-null	float64
6	v6	284807	non-null	float64
7	v7	284807	non-null	float64
8	v8	284807	non-null	float64
9	v9	284807	non-null	float64
10	v10	284807	non-null	float64
11	v11	284807	non-null	float64
12	v12	284807	non-null	float64
13	v13	284807	non-null	float64
14	v14	284807	non-null	float64
15	v15	284807	non-null	float64
16	v16	284807	non-null	float64
17	v17	284807	non-null	float64
18	v18	284807	non-null	float64
19	v19	284807	non-null	float64
20	v20	284807	non-null	float64
21	v21	284807	non-null	float64
22	v22	284807	non-null	float64
23	v23	284807	non-null	float64
24	v24	284807	non-null	float64
25	v25	284807	non-null	float64
26	v26	284807	non-null	float64
27	v27	284807	non-null	float64
28	v28	284807	non-null	float64
29	amount	284807	non-null	float64
30	class	284807	non-null	int64

dtypes: float64(30), int64(1)

memory usage: 67.4 MB

None

```
In [6]: # List all columns in the dataset.  
print("Columns in Dataset")  
print(df.columns)
```

Columns in Dataset

```
Index(['time', 'v1', 'v2', 'v3', 'v4', 'v5', 'v6', 'v7', 'v8', 'v9', 'v10',  
      'v11', 'v12', 'v13', 'v14', 'v15', 'v16', 'v17', 'v18', 'v19', 'v20',  
      'v21', 'v22', 'v23', 'v24', 'v25', 'v26', 'v27', 'v28', 'amount',  
      'class'],  
      dtype='object')
```

```
In [7]: print("\nData shape:", df.shape)
```

Data shape: (284807, 31)

```
In [8]: # Count and print missing values per column  
print(df.isnull().sum())
```

```
time      0
v1         0
v2         0
v3         0
v4         0
v5         0
v6         0
v7         0
v8         0
v9         0
v10        0
v11        0
v12        0
v13        0
v14        0
v15        0
v16        0
v17        0
v18        0
v19        0
v20        0
v21        0
v22        0
v23        0
v24        0
v25        0
v26        0
v27        0
v28        0
amount     0
class      0
dtype: int64
```

Data Cleaning

```
In [9]: # Check unique values & class balance

print("\nUnique values in 'Class':", df['class'].unique())
print("\nClass distribution (counts):")
```

```
print(df['class'].value_counts())
print("\nClass distribution (percentage):")
print(df['class'].value_counts(normalize=True) * 100)
```

Unique values in 'Class': [0 1]

Class distribution (counts):

class

0 284315

1 492

Name: count, dtype: int64

Class distribution (percentage):

class

0 99.827251

1 0.172749

Name: proportion, dtype: float64

RESULT

So there are 284,315 normal transactions and 492 fraud transaction. The dataset is heavily imbalanced.

```
In [44]: # Simple textual report
# -----
fraud_count = df['class'].value_counts()[1]
normal_count = df['class'].value_counts()[0]

print("Quick Report Summary:")
print(f"Total transactions: {len(df)}")
print(f"Normal transactions: {normal_count}")
print(f"Fraud transactions: {fraud_count}")
print(f"Fraud Rate: {100 * fraud_count / len(df):.4f}%")

print("\nData is ready for modeling or deeper insights ")
```

Quick Report Summary:
Total transactions: 284807
Normal transactions: 284315
Fraud transactions: 492
Fraud Rate: 0.1727%

Data is ready for modeling or deeper insights

Data Validation

```
In [45]: # Data validation on SQL and Python
df = pd.read_sql("SELECT COUNT(*) FROM cc_data", engine)
print("SQL:", df)
print(f"✅ Row count validation passed: {df.shape[0]} rows match SQL.")
```

```
SQL:      count
0  284807
✅ Row count validation passed: 1 rows match SQL.
```

Report Genration

1. Dataset Identification ✅
2. Data Import into SQL ✅
3. Normalization into relational tables : our table is flat table. No normalization. As data didn't require normalization as each row is an independent transaction. ✅
4. Initial SQL Profiling (SQL+Python) ✅
5. Data Cleaning & Transformation (Python): Nulls (none found), Duplicates (kept fraud class and non-fraud duplicates), Cleaning on Task 3 ✅
6. Data Validation ✅
7. Final Deliverable : Cleaned dataset on task 3 ✅

EXPLORATORY DATA ANALYSIS AND VISUALIZATION

Summary Statistics

```
In [10]: # Exploratory data analysis  
  
print("\nStatistical summary:")  
print(df.describe())
```


Statistical summary:

	time	v1	v2	v3	v4 \
count	284807.000000	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	94813.859575	1.168375e-15	3.416908e-16	-1.379537e-15	2.074095e-15
std	47488.145955	1.958696e+00	1.651309e+00	1.516255e+00	1.415869e+00
min	0.000000	-5.640751e+01	-7.271573e+01	-4.832559e+01	-5.683171e+00
25%	54201.500000	-9.203734e-01	-5.985499e-01	-8.903648e-01	-8.486401e-01
50%	84692.000000	1.810880e-02	6.548556e-02	1.798463e-01	-1.984653e-02
75%	139320.500000	1.315642e+00	8.037239e-01	1.027196e+00	7.433413e-01
max	172792.000000	2.454930e+00	2.205773e+01	9.382558e+00	1.687534e+01

	v5	v6	v7	v8	v9 \
count	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	9.604066e-16	1.487313e-15	-5.556467e-16	1.205498e-16	-2.406306e-15
std	1.380247e+00	1.332271e+00	1.237094e+00	1.194353e+00	1.098632e+00
min	-1.137433e+02	-2.616051e+01	-4.355724e+01	-7.321672e+01	-1.343407e+01
25%	-6.915971e-01	-7.682956e-01	-5.540759e-01	-2.086297e-01	-6.430976e-01
50%	-5.433583e-02	-2.741871e-01	4.010308e-02	2.235804e-02	-5.142873e-02
75%	6.119264e-01	3.985649e-01	5.704361e-01	3.273459e-01	5.971390e-01
max	3.480167e+01	7.330163e+01	1.205895e+02	2.000721e+01	1.559499e+01

	...	v21	v22	v23	v24 \
count	...	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	...	1.656562e-16	-3.568593e-16	2.610582e-16	4.473066e-15
std	...	7.345240e-01	7.257016e-01	6.244603e-01	6.056471e-01
min	...	-3.483038e+01	-1.093314e+01	-4.480774e+01	-2.836627e+00
25%	...	-2.283949e-01	-5.423504e-01	-1.618463e-01	-3.545861e-01
50%	...	-2.945017e-02	6.781943e-03	-1.119293e-02	4.097606e-02
75%	...	1.863772e-01	5.285536e-01	1.476421e-01	4.395266e-01
max	...	2.720284e+01	1.050309e+01	2.252841e+01	4.584549e+00

	v25	v26	v27	v28	amount \
count	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	284807.000000
mean	5.213180e-16	1.683537e-15	-3.659966e-16	-1.223710e-16	88.349619
std	5.212781e-01	4.822270e-01	4.036325e-01	3.300833e-01	250.120109
min	-1.029540e+01	-2.604551e+00	-2.256568e+01	-1.543008e+01	0.000000
25%	-3.171451e-01	-3.269839e-01	-7.083953e-02	-5.295979e-02	5.600000
50%	1.659350e-02	-5.213911e-02	1.342146e-03	1.124383e-02	22.000000
75%	3.507156e-01	2.409522e-01	9.104512e-02	7.827995e-02	77.165000
max	7.519589e+00	3.517346e+00	3.161220e+01	3.384781e+01	25691.160000

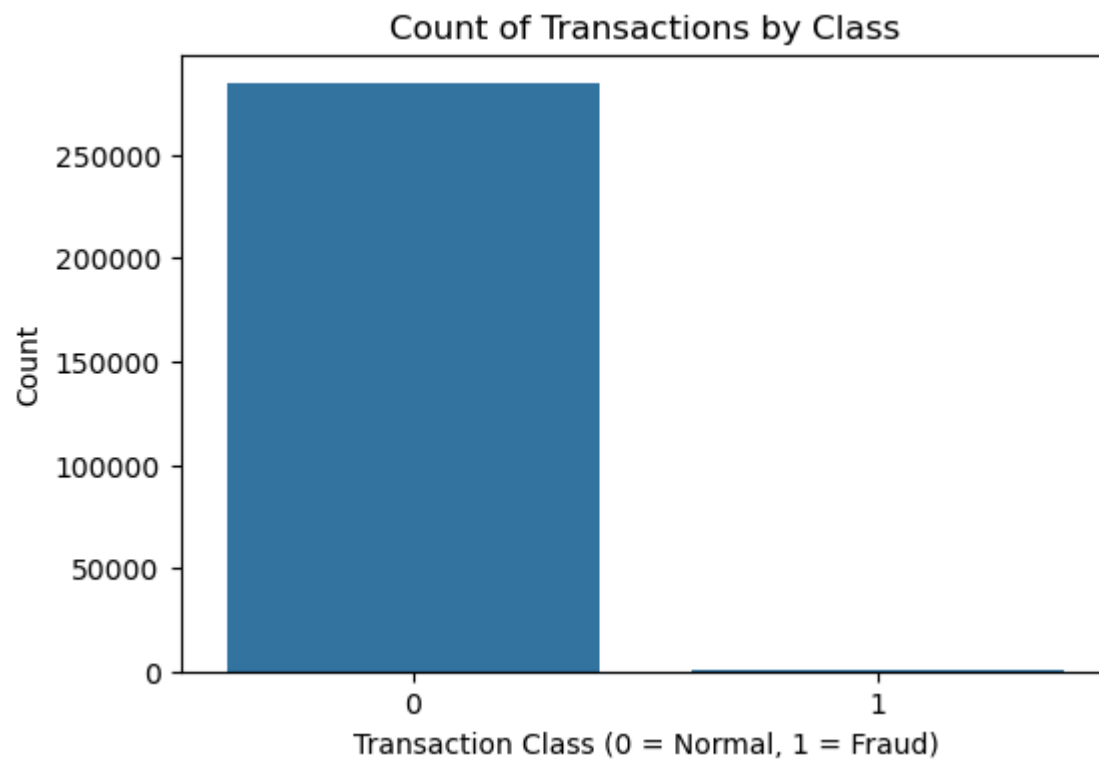
```
      class
count  284807.000000
mean    0.001727
std     0.041527
min     0.000000
25%    0.000000
50%    0.000000
75%    0.000000
max     1.000000
```

```
[8 rows x 31 columns]
```

Univariate Analysis

```
In [11]: # Distribution of normal transaction and fraud transaction
```

```
plt.figure(figsize=(6, 4))
sns.countplot(x='class', data=df)
plt.title('Count of Transactions by Class')
plt.xlabel('Transaction Class (0 = Normal, 1 = Fraud)')
plt.ylabel('Count')
plt.show()
```



```
In [12]: # Count and print unique transaction amounts  
print(df["amount"].value_counts())
```

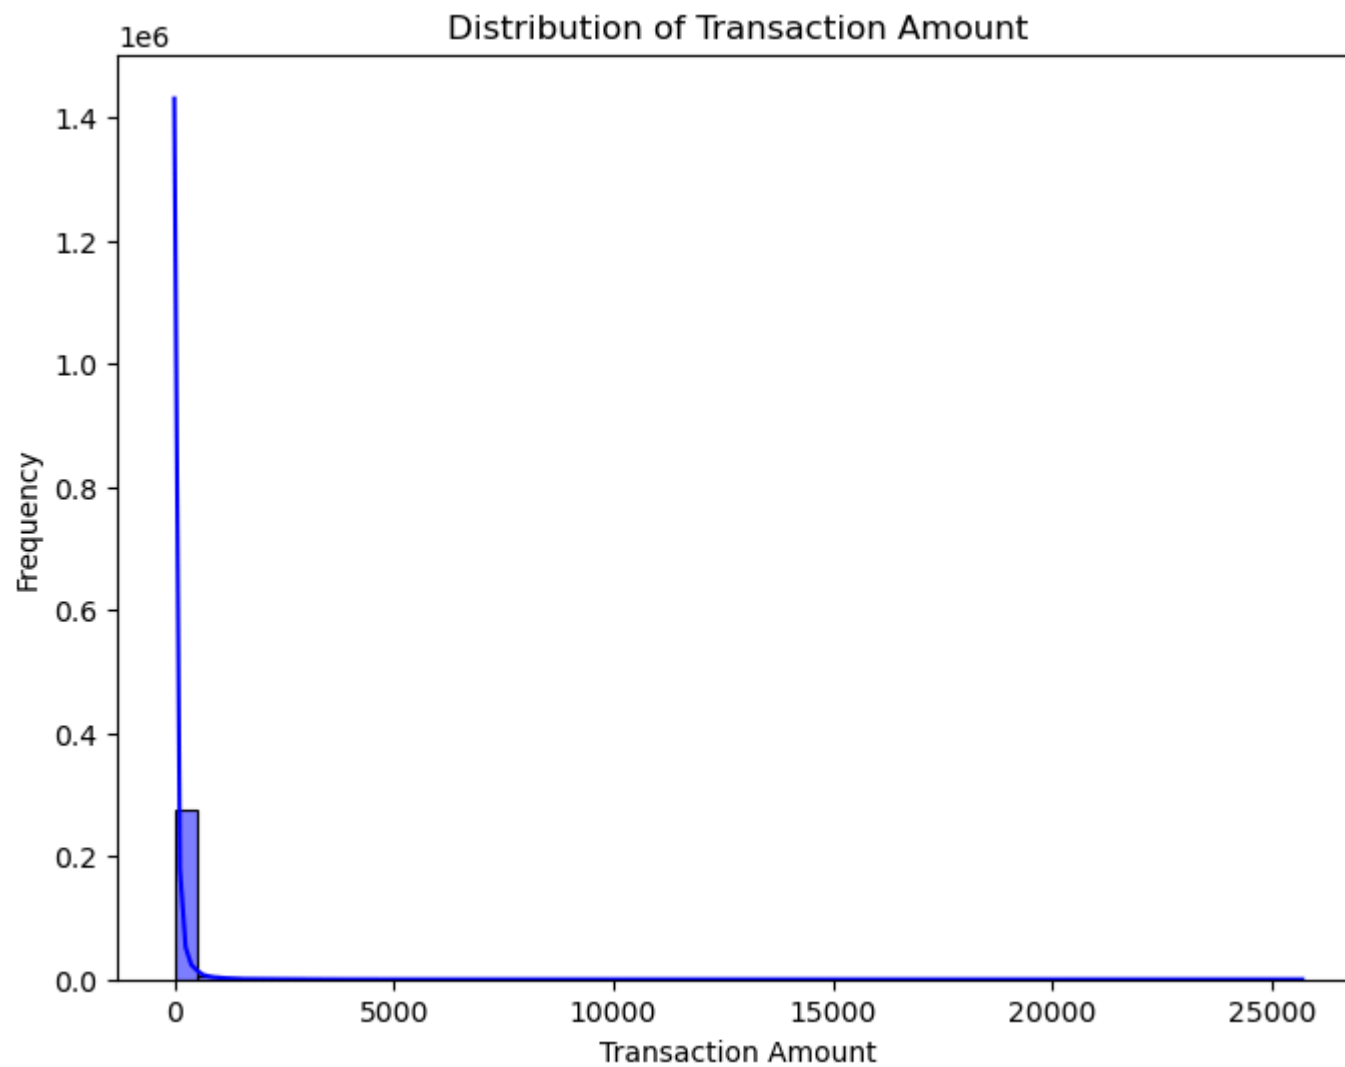
```
amount  
1.00      13688  
1.98       6044  
0.89       4872  
9.99       4747  
15.00      3280  
...  
202.24         1  
252.85         1  
615.52         1  
180.93         1  
807.48         1  
Name: count, Length: 32767, dtype: int64
```

```
In [13]: # Show percentage distribution of transaction amounts
df["amount"].value_counts(normalize=True)
```

```
Out[13]: amount
1.00      0.048061
1.98      0.021221
0.89      0.017106
9.99      0.016667
15.00     0.011517
...
202.24    0.000004
252.85    0.000004
615.52    0.000004
180.93    0.000004
807.48    0.000004
Name: proportion, Length: 32767, dtype: float64
```

```
In [14]: # Distribution of transaction Amount

plt.figure(figsize=(8, 6))
sns.histplot(df['amount'], bins=50, kde=True, color='blue')
plt.title('Distribution of Transaction Amount')
plt.xlabel('Transaction Amount')
plt.ylabel('Frequency')
plt.show()
```



```
In [15]: # Count and print unique transaction times (duration)
print(df["time"].value_counts())
```

```
time
163152.0    36
64947.0     26
68780.0     25
3767.0      21
3770.0      20
..
172760.0     1
172758.0     1
172757.0     1
172756.0     1
172754.0     1
Name: count, Length: 124592, dtype: int64
```

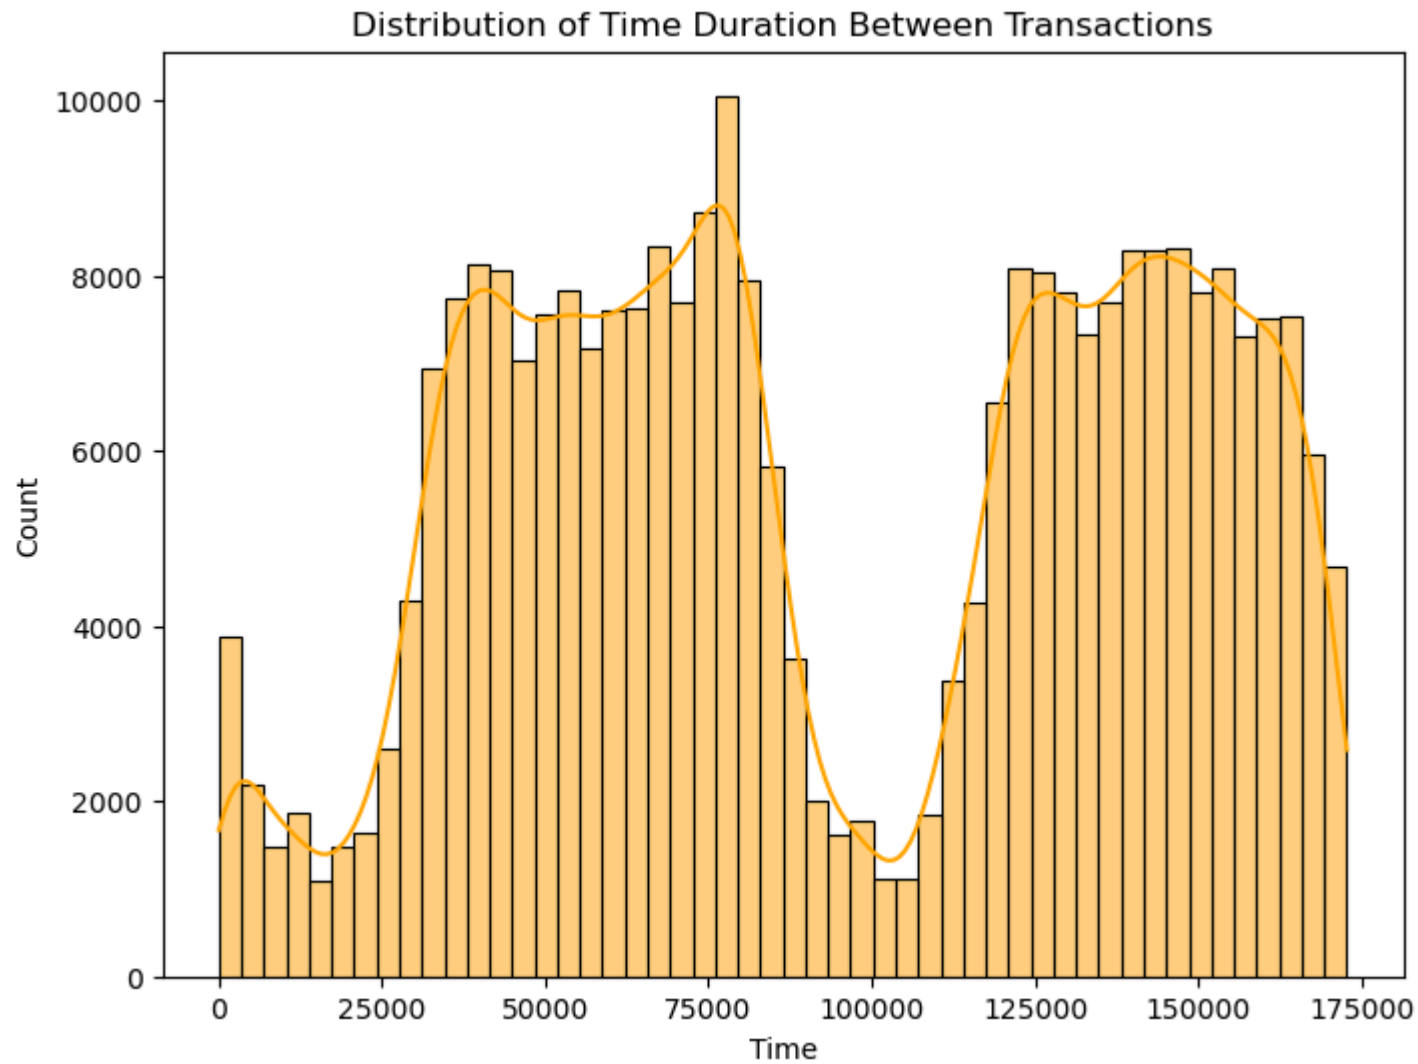
```
In [16]: # Show percentage distribution of transaction times (duration)
df["time"].value_counts(normalize=True)
```

```
Out[16]: time
163152.0    0.000126
64947.0     0.000091
68780.0     0.000088
3767.0      0.000074
3770.0      0.000070
...
172760.0    0.000004
172758.0    0.000004
172757.0    0.000004
172756.0    0.000004
172754.0    0.000004
Name: proportion, Length: 124592, dtype: float64
```

```
In [17]: # Distribution of Time duration between transactions

plt.figure(figsize=(8,6))
sns.histplot(df["time"], bins=50, kde=True, color="orange")
plt.title("Distribution of Time Duration Between Transactions")
plt.xlabel("Time")
plt.ylabel("Count")
```

```
Out[17]: Text(0, 0.5, 'Count')
```



Bivariate/ Multivariate Analysis

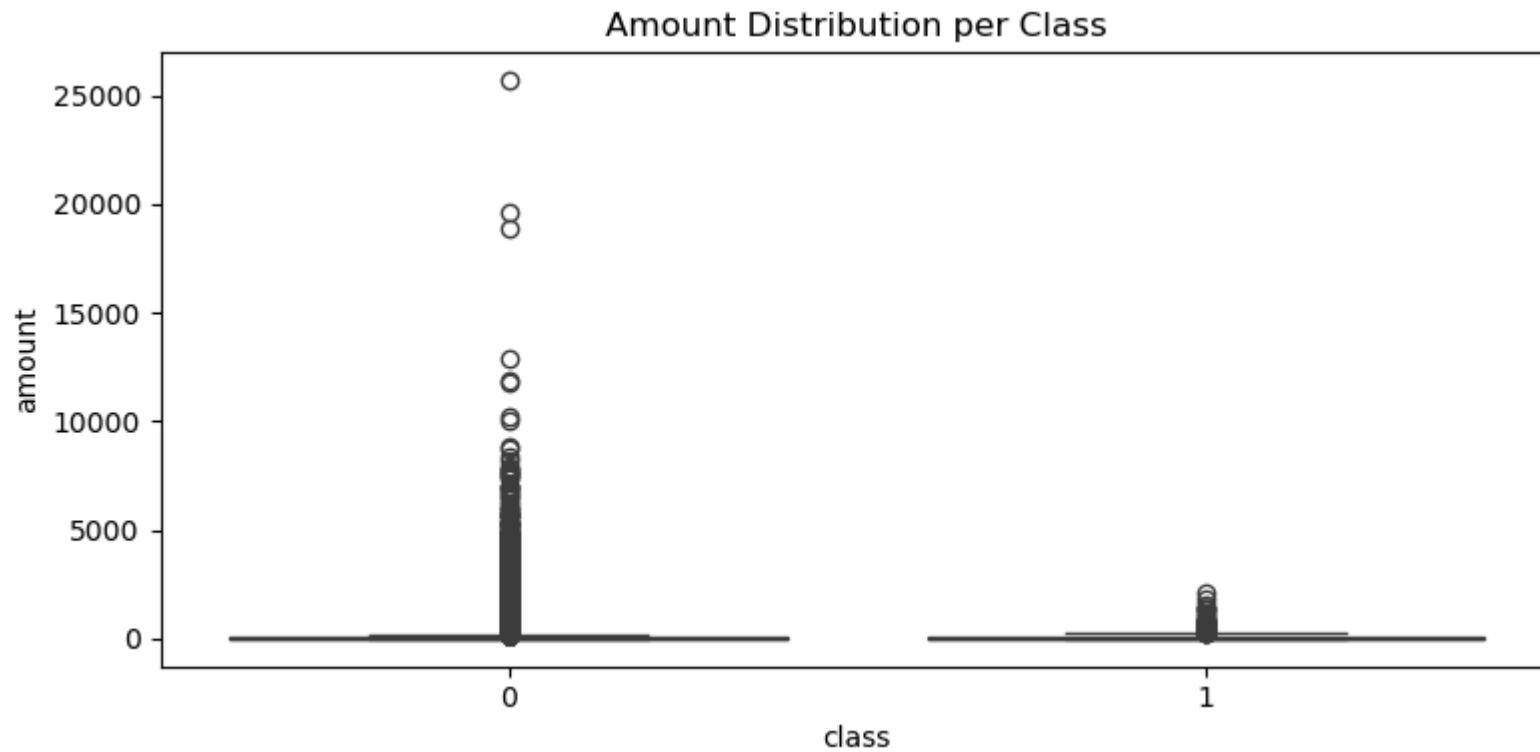
```
In [18]: # Distribution of Time and Amount per Class (0=non-fraud, 1=fraud case)

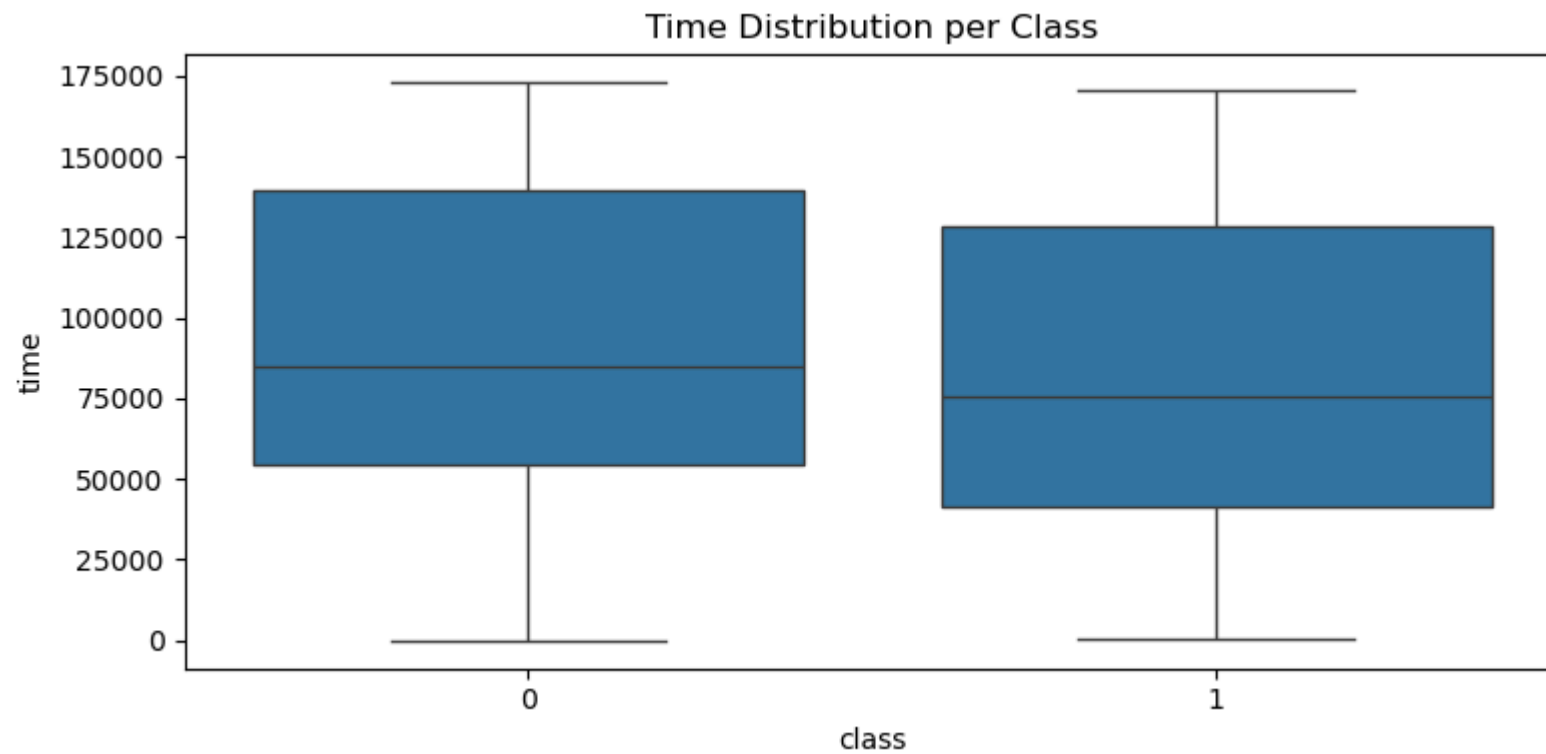
# Boxplot Figure 1: Amount vs Class
```

```
plt.figure(figsize=(8,4))
sns.boxplot(x="class", y="amount", data=df)
plt.title("Amount Distribution per Class")
plt.tight_layout()
plt.show()
```

Boxplot Figure 2: Time vs Class

```
plt.figure(figsize=(8,4))
sns.boxplot(x="class", y="time", data=df)
plt.title("Time Distribution per Class")
plt.tight_layout()
plt.show()
```





```
In [19]: # Correlation between Features and Class 1 (Fraud)

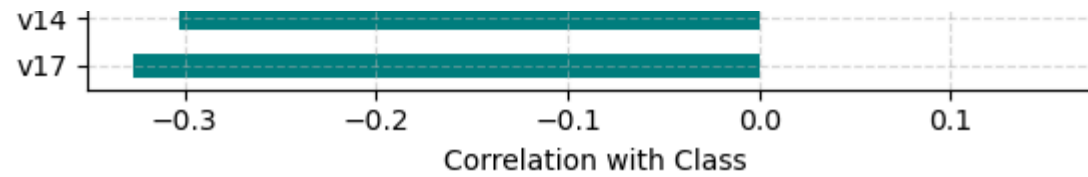
correlations = df.corr()["class"].drop("class").sort_values()

plt.figure(figsize=(6, 8))
correlations.plot(kind="barh", color="teal")

plt.title("Other Features Correlation with Class (FRAUD)")
plt.xlabel("Correlation with Class")
plt.grid(True, linestyle='--', alpha=0.5)
plt.tight_layout()
plt.show()
```

Other Features Correlation with Class (FRAUD)





DATA PREPARATION

```
In [20]: # 1. Handling missing values.  
df.isnull().sum()
```

```
Out[20]: time      0
         v1        0
         v2        0
         v3        0
         v4        0
         v5        0
         v6        0
         v7        0
         v8        0
         v9        0
        v10        0
        v11        0
        v12        0
        v13        0
        v14        0
        v15        0
        v16        0
        v17        0
        v18        0
        v19        0
        v20        0
        v21        0
        v22        0
        v23        0
        v24        0
        v25        0
        v26        0
        v27        0
        v28        0
        amount     0
        class      0
        dtype: int64
```

```
In [21]: # 2.Handling duplicated values.
```

```
# Show all rows that are duplicates (keeping and showing all occurrences)
df[df.duplicated(keep=False)]
```

Out[21]:

	time	v1	v2	v3	v4	v5	v6	v7	v8	v9	...	v21	v22
32	26.0	-0.529912	0.873892	1.347247	0.145457	0.414209	0.100223	0.711206	0.176066	-0.286717	...	0.046949	0.208105
33	26.0	-0.529912	0.873892	1.347247	0.145457	0.414209	0.100223	0.711206	0.176066	-0.286717	...	0.046949	0.208105
34	26.0	-0.535388	0.865268	1.351076	0.147575	0.433680	0.086983	0.693039	0.179742	-0.285642	...	0.049526	0.206537
35	26.0	-0.535388	0.865268	1.351076	0.147575	0.433680	0.086983	0.693039	0.179742	-0.285642	...	0.049526	0.206537
112	74.0	1.038370	0.127486	0.184456	1.109950	0.441699	0.945283	-0.036715	0.350995	0.118950	...	0.102520	0.605089
...
283485	171627.0	-1.457978	1.378203	0.811515	-0.603760	-0.711883	-0.471672	-0.282535	0.880654	0.052808	...	0.284205	0.949659
284190	172233.0	-2.667936	3.160505	-3.355984	1.007845	-0.377397	-0.109730	-0.667233	2.309700	-1.639306	...	0.391483	0.266536
284191	172233.0	-2.667936	3.160505	-3.355984	1.007845	-0.377397	-0.109730	-0.667233	2.309700	-1.639306	...	0.391483	0.266536
284192	172233.0	-2.691642	3.123168	-3.339407	1.017018	-0.293095	-0.167054	-0.745886	2.325616	-1.634651	...	0.402639	0.259746
284193	172233.0	-2.691642	3.123168	-3.339407	1.017018	-0.293095	-0.167054	-0.745886	2.325616	-1.634651	...	0.402639	0.259746

1854 rows × 31 columns



```
In [22]: # Show duplicates based only on Time, Amount, and Class features
df[df.duplicated(subset=["time", "amount", "class"], keep=False)]
```

Out[22]:

	time	v1	v2	v3	v4	v5	v6	v7	v8	v9	...	v21	v22
32	26.0	-0.529912	0.873892	1.347247	0.145457	0.414209	0.100223	0.711206	0.176066	-0.286717	...	0.046949	0.208109
33	26.0	-0.529912	0.873892	1.347247	0.145457	0.414209	0.100223	0.711206	0.176066	-0.286717	...	0.046949	0.208109
34	26.0	-0.535388	0.865268	1.351076	0.147575	0.433680	0.086983	0.693039	0.179742	-0.285642	...	0.049526	0.206537
35	26.0	-0.535388	0.865268	1.351076	0.147575	0.433680	0.086983	0.693039	0.179742	-0.285642	...	0.049526	0.206537
108	73.0	1.162281	1.248178	-1.581317	1.475024	1.138357	-1.020373	0.638387	-0.136762	-0.805505	...	-0.124012	-0.227150
...
284193	172233.0	-2.691642	3.123168	-3.339407	1.017018	-0.293095	-0.167054	-0.745886	2.325616	-1.634651	...	0.402639	0.259740
284248	172273.0	-0.765414	1.343887	-0.306101	-0.645545	-0.067358	-1.172196	0.516073	0.342927	0.368227	...	-0.289752	-0.709880
284251	172273.0	2.061056	-0.077031	-1.068720	0.422266	-0.181192	-1.227747	0.160285	-0.314824	0.596385	...	-0.292782	-0.727550
284328	172348.0	2.064806	0.008284	-2.226901	0.926502	1.119908	0.178604	0.349210	-0.010441	0.262333	...	0.006761	0.087820
284329	172348.0	-1.351689	1.969541	-2.145252	-0.866654	0.438384	-0.124297	-0.245481	1.404284	-0.342847	...	-0.296305	-1.007130

8736 rows × 31 columns



DATA CLEANING

```
In [23]: # Count the total number of duplicated rows (keeping the first occurrence as non-duplicate)  
df.duplicated().sum()
```

Out[23]: np.int64(1081)

```
In [24]: # Count the total number of duplicated rows (marking all duplicates as True)  
df.duplicated(keep=False).sum()
```

Out[24]: np.int64(1854)

```
In [25]: # Count the number of unique duplicated rows (after dropping exact duplicates)
df[df.duplicated(keep=False)].drop_duplicates().shape[0]
```

```
Out[25]: 773
```

```
In [26]: # Show the distribution of Class values among duplicated rows
df[df.duplicated(keep=False)]["class"].value_counts()
```

```
Out[26]: class
0      1822
1         32
Name: count, dtype: int64
```

Observation

There are 1,822 duplicats on Class 0 (normal) and 32 duplicats on Class 1 (fraud).

Since Class 1 has few duplicats, I decided not too drop the duplicats.

I will only drop the duplicats from Class 0.

```
In [30]: # Separating Class between 0 and 1
fraud = df[df["class"] == 1]
normal = df[df["class"] == 0]

# Drop duplicated values from Class 0 only
normal_cleaned = normal.drop_duplicates()

# Combined both Classes and renamed the dataset as df_cleaned

# Duplicated values of Class 0 have been dropped from the dataset
df_cleaned = pd.concat([fraud, normal_cleaned], ignore_index=True)
```

```
In [31]: # Count duplicates in the cleaned dataset (df_cleaned)
df_cleaned.duplicated().sum()
```

```
Out[31]: np.int64(19)
```

```
In [32]: # Check Class distribution in duplicates of the cleaned dataset
df_cleaned[df_cleaned.duplicated(keep=False)]["class"].value_counts()
```

```
Out[32]: class
1      32
Name: count, dtype: int64
```

```
In [33]: # Show the percentage distribution of unique values in the Class feature
df_cleaned["class"].value_counts(normalize=True) * 100
```

```
Out[33]: class
0      99.826605
1       0.173395
Name: proportion, dtype: float64
```

Outliers and Anomaly Detection

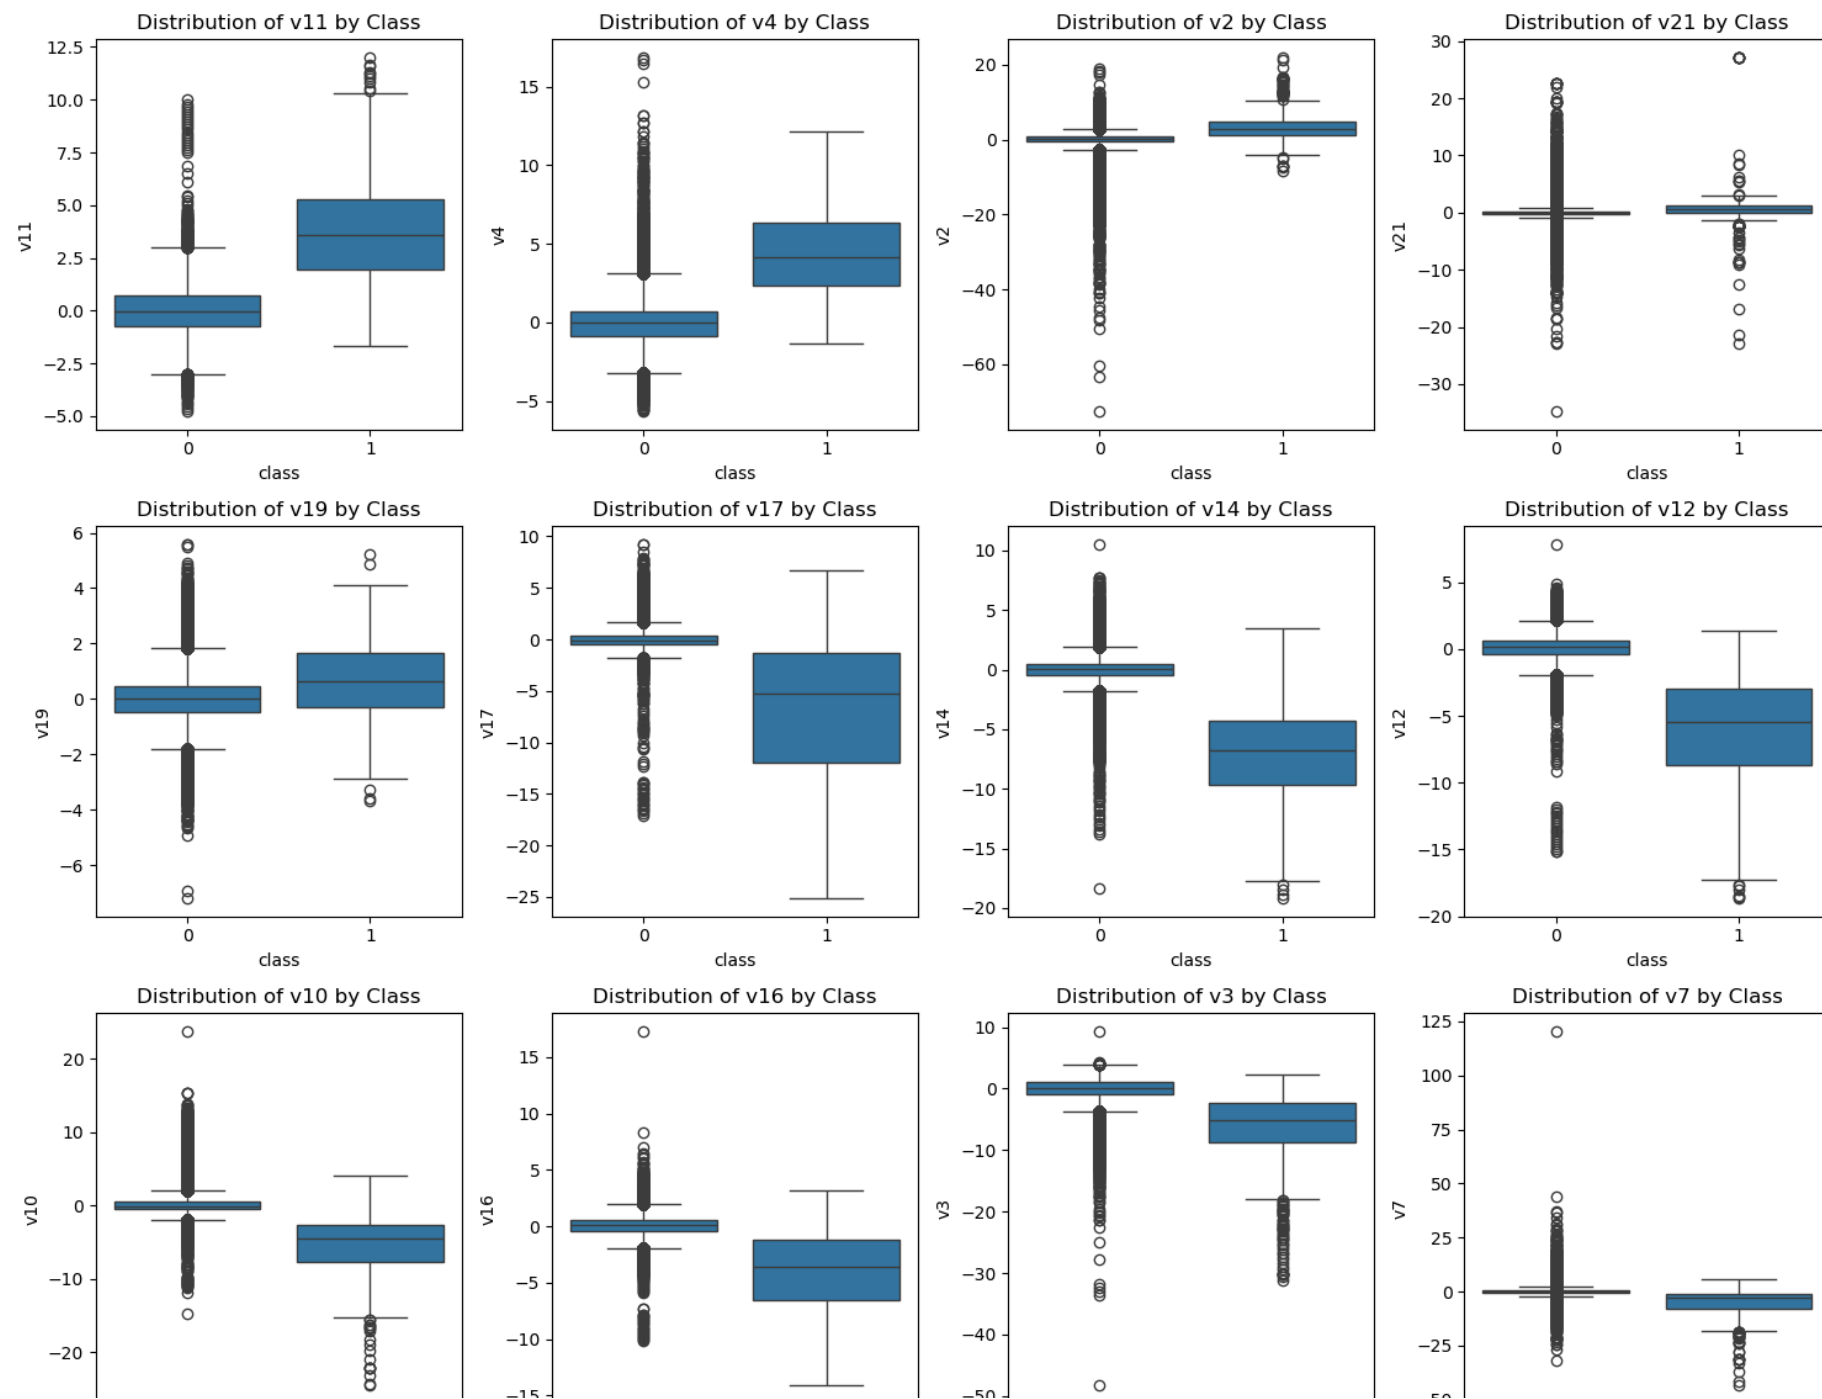
```
In [39]: #CHECKING CLASS DISTRIBUTION:

# A. Feature Distribution between Fraud and Non-fraud Transactions

# Selected features based on Strongest Positive and Strongest Negative correlations with Class

selected_features = ['v11', 'v4', 'v2', 'v21', 'v19', # Strong Positive
                    'v17', 'v14', 'v12', 'v10', 'v16', 'v3', 'v7'] # Strong Negative

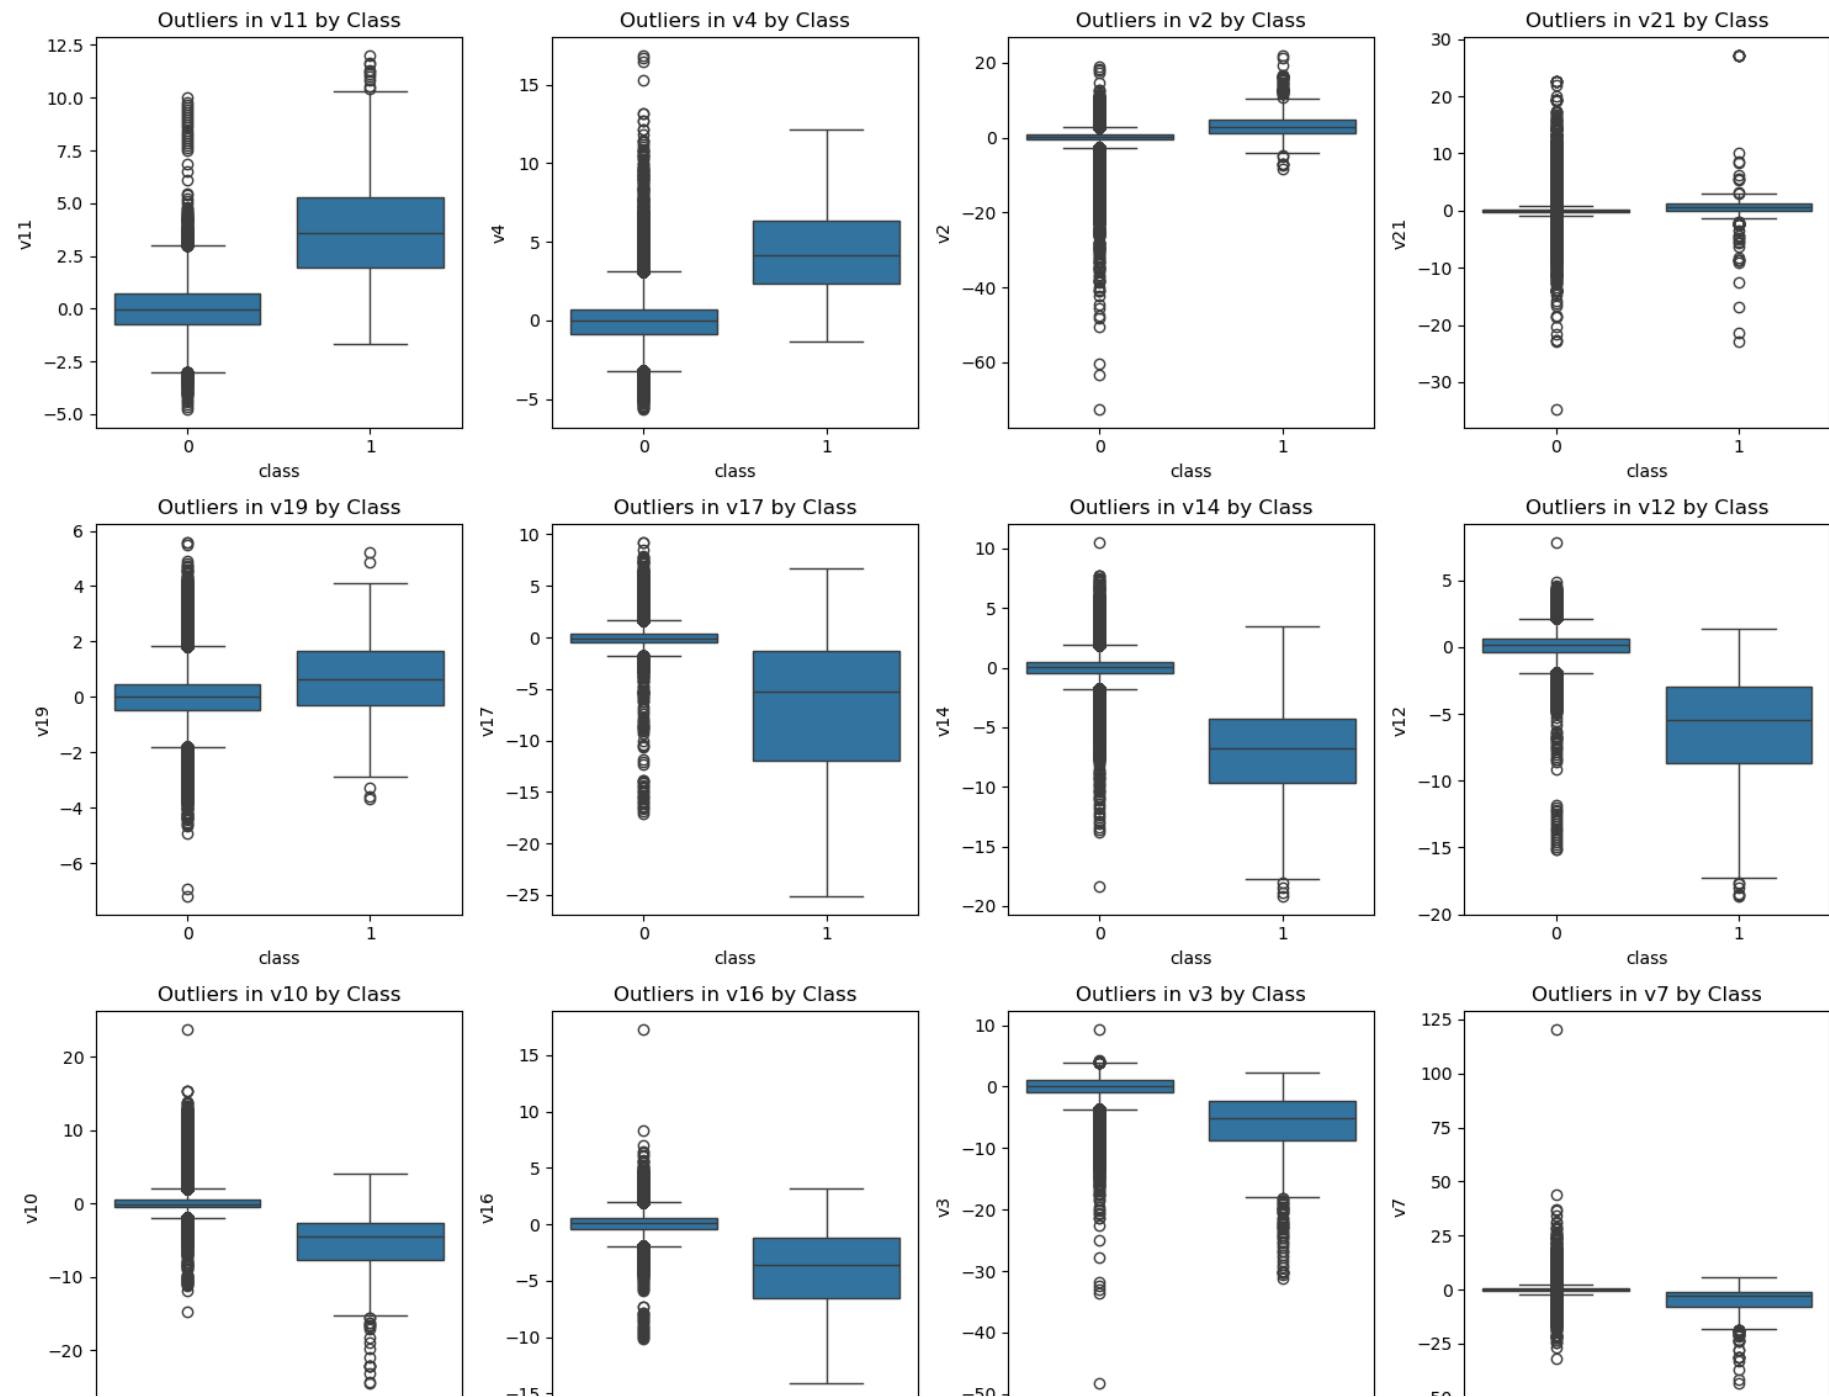
plt.figure(figsize=(15, 12))
for i, feature in enumerate(selected_features, 1):
    plt.subplot(3, 4, i)
    sns.boxplot(data=df_cleaned, x='class', y=feature)
    plt.title(f'Distribution of {feature} by Class')
plt.tight_layout()
plt.show()
```



```
In [40]: # B. Outlier Visualization with Boxplots

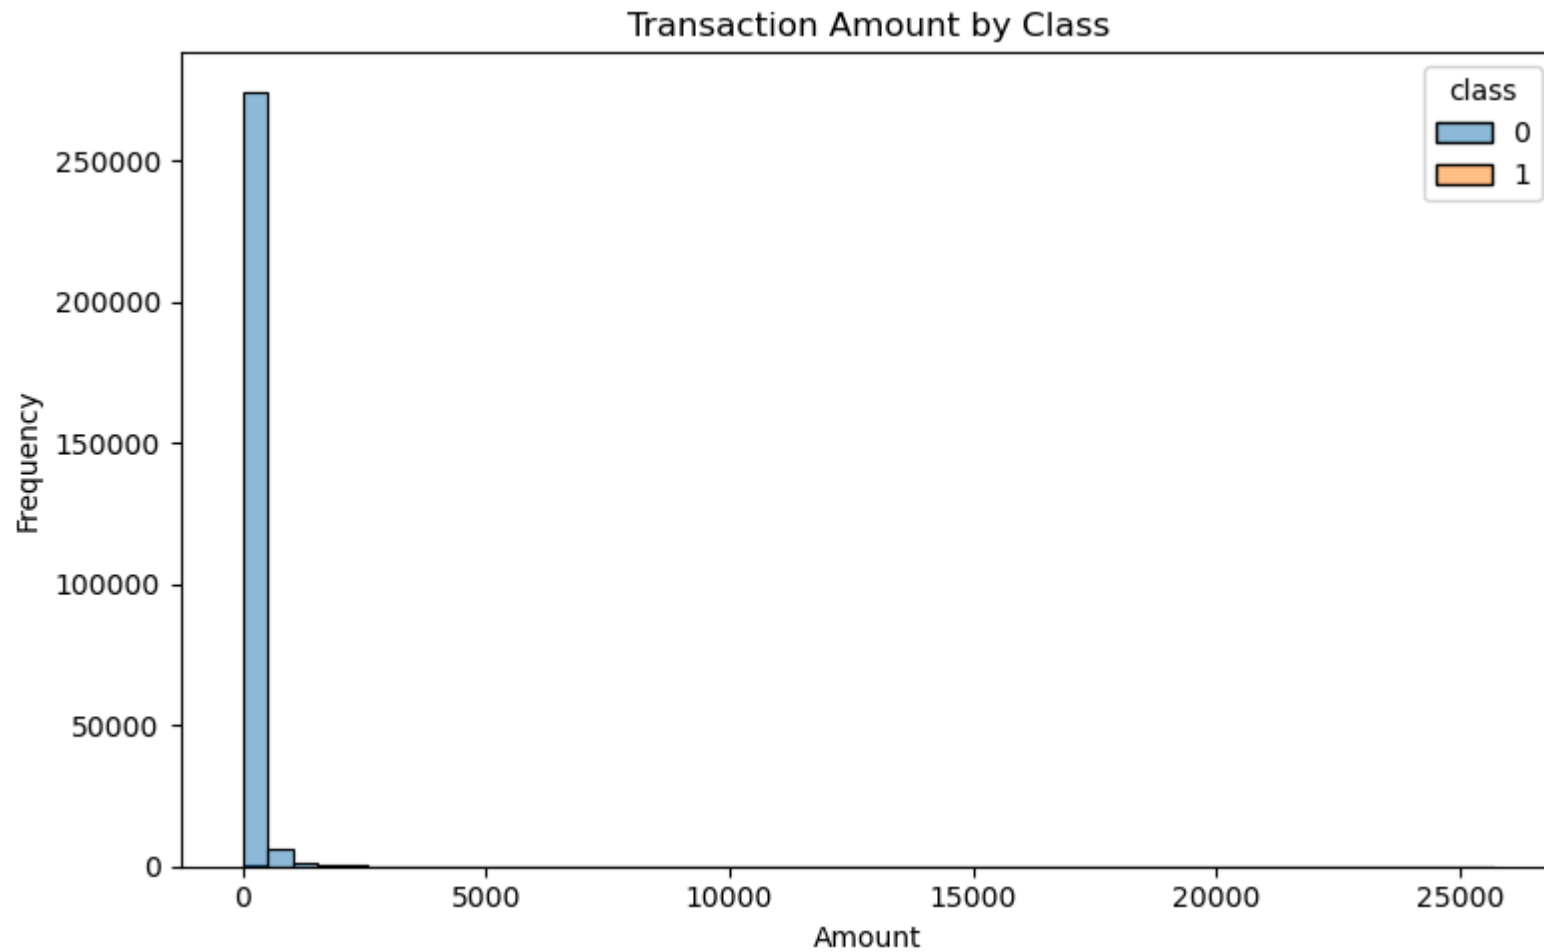
plt.figure(figsize=(15, 12))
for i, feature in enumerate(selected_features, 1):
    plt.subplot(3, 4, i)
    sns.boxplot(data=df_cleaned, x='class', y=feature)
    plt.title(f'Outliers in {feature} by Class')
plt.tight_layout()
plt.show()
```





In [41]: *# C. Histogram for Transaction Amount*

```
plt.figure(figsize=(8, 5))
sns.histplot(df_cleaned, x='amount', hue='class', bins=50)
plt.title('Transaction Amount by Class', fontsize=12)
plt.xlabel('Amount')
plt.ylabel('Frequency')
plt.tight_layout()
plt.show()
```



RESULT

- I have decided not to drop any outliers.
- Outliers were retained because they may represent real fraudulent behavior, which is inherently anomalous.
- Removing them could reduce the model's ability to detect rare but significant fraud patterns.

PowerBI Visulization

```
In [42]: # Find all duplicates (marking all occurrences, not just subsequent ones)
duplicates_df = df[df.duplicated(keep=False)]

# KPIs for duplicates
total_duplicates = duplicates_df.shape[0]          # ALL duplicates found
duplicate_class_counts = duplicates_df['class'].value_counts()

# Breakdown
duplicate_normal = duplicate_class_counts.get(0, 0)
duplicate_fraud = duplicate_class_counts.get(1, 0)

# After cleaning
total_after_cleaning = df_cleaned.shape[0]

print(f"Total duplicate records found: {total_duplicates}")
print(f"Duplicates in Normal class: {duplicate_normal}")
print(f"Duplicates in Fraud class: {duplicate_fraud}")
print(f"Total records after cleaning: {total_after_cleaning}")
```

```
Total duplicate records found: 1854
Duplicates in Normal class: 1822
Duplicates in Fraud class: 32
Total records after cleaning: 283745
```

```
In [43]: # Export to CSV
df_cleaned.to_csv("cleaned_creditcard_data.csv", index=False)
```

POWER BI DATA VISUALIZATION

- Total Transactions - 284,807
- Total Fraud Cases - 492
- Fraud Rate (%) - 0.1727%
- Total Normal Transactions - 284,315
- Average Transaction Amount - 88.35 (currency units)

- Max Transaction Amount -25,691.16
- Min Transaction Amount - 0.00
- ☒ No missing values in dataset.
- ☒ Duplicates found:
- Total Duplicate Records - 1854
- Duplicates in Normal Class - 1822
- Duplicates in Fraud Class - 32
- Total Records After Cleaning - 283746

Final Clean Dataset

- Total Records (after cleaning) 283746
- Fraud Rate (after cleaning) ~0.17% unchanged